

# Fashion-MNIST Final Classification Project

DONGHEE LEE, VICTOR SUCIU, and LINH HA

## 1 INTRODUCTION

The following report explores the usage of artificial neural networks to classify Fashion-MNIST. Fashion-MNIST is a data set of 60,000 grey-scaled clothing images. For our first set of experiments, we embarked on a grand journey of trial-and-error and tested various configurations of a multilayered perceptron that utilized backpropagation. We modified hyper-parameters, changed the number of layers and neurons, and adjusted the transfer functions. Next, we augmented the training data by adding multiple slightly altered copies of each image. Finally, we created an Ensemble model that aggregated multiple trained MLPs into a single, more intelligent model. We tracked the accuracy of each model and produced a csv submission file for the *UWB CSS 485: Winter 2021 Project* Kaggle competition.

## 2 EXPERIMENTS - METHODS

First, we started training a wide variety of MLPs in an attempt to increase classification accuracy. We experimented with the number of layers, neuron counts, and hyperparameters for gradient descent. After we achieved reasonably high performance, we implemented some data augmentation techniques to increase the size of the training set and experimented further with MLP configurations. *Appendix 7.1 Multilayer Perceptron Class* and *Appendix 7.2 Perceptron Layer Class* are the classes used to build our network. *Appendix 7.6 Transfer Functions and their Derivatives* are the different functions we used for our network.

### 2.1 Preprocessing, Training, and Prediction

Before training, we performed some minor preprocessing to the data. First, the numeric labels were converted to one-hot format using *Appendix 7.11 One-Hot Encoding and Decoding* and the images were reshaped into 784-dimensional column vectors. Next, each image was normalized by dividing each element by 255 which reduced their range to  $[0, 1]$ . Then, we divided our data set into a training set with 50,000 data points and a validation set of 10,000 data points.

Training was done with a batch size ranging from 24 to 64 and a variety of epoch size. For epochs, we did not use a fixed number automatic early stopping. Instead, MultilayerPerceptron's training function checked the validation accuracy after every epoch, and saved a checkpoint if the accuracy was a new maximum. Each checkpoint was an entire MultilayerPerceptron object that was serialized as a file. This meant that any checkpoint could be loaded back into MatLab for prediction or further training. This strategy allowed us to continue training a model well into overfitting territory while leaving the best version untarnished. *Appendix 7.12 Main Function - Standard* demonstrates the driver code that builds and trains the model.

Once a model was trained, we loaded in its best checkpoint and used it to predict the submission data. Final prediction was done by passing data through the model's forward function and applying the hardmax function to the output. Hardmax, also known as argmax, takes a vector as input and sets maximum element to 1 and the rest to 0. This results in the same one-hot format as the labels where the maximum vector element is the model's prediction. *Appendix 7.8 Hardmax Function* shows the implementation of hardmax.

Table 1. All Unique MLP Configurations with Best Accuracy Per-Configuration

Layer 1		Layer 2		Layer 3		Layer 4		Params			Acc %
Trans	Neur	Trans	Neur	Trans	Neur	Trans	Neur	Lr	Mtm	Dcy	
50,000 Train — 10,000 Validation											
tanh	256	tanh	128	tanh	48	sigmoid	10	0.1	0	0	68.00
linear	256	linear	128	linear	48	sigmoid	10	0.05	0	0	84.00
sigmoid	400	sigmoid	250	sigmoid	100	relu	10	0.1	0	0	86.86
sigmoid	400	sigmoid	250	sigmoid	100	softmax	10	0.07	0.8	0	88.40
sigmoid	400	sigmoid	250	sigmoid	100	relu	10	0.08	0.5	0	88.40
sigmoid	400	sigmoid	250	sigmoid	100	relu	10	0.05	0.8	0	89.10
sigmoid	400	sigmoid	250	sigmoid	100	relu	10	0.03	1	0	89.24
250,000 Augmented Train — 10,000 Validation											
sigmoid	400	sigmoid	250	sigmoid	100	relu	10	0.03	1	0	89.40
sigmoid	400	sigmoid	250	sigmoid	100	relu	10	0.07	1	0	89.63
tanh	400	sigmoid	250	sigmoid	100	softmax	10	0.05	0.9	0	89.26
sigmoid	400	sigmoid	150	softmax	10	—	—	0.08	1	0	90.14
sigmoid	400	sigmoid	150	softmax	10	—	—	0.08	1	-0.01	90.40
300,000 Augmented Train — 10,000 Validation											
sigmoid	500	sigmoid	250	relu	10	—	—	0.08	1	-0.01	89.92

Every unique multilayer perceptron configuration we tried and the best validation accuracy for each one. Many of these were trained multiple times to account for random weight initialization, but only the best accuracy is shown for each configuration.

Key:

Trans - Transfer Function   Neur - Number of Neurons   Lr - Learning Rate   Mtm - Momentum   Dcy - Weight Decay

Acc - Validation Accuracy

## 2.2 Overview of Trials

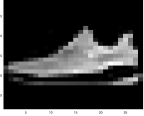
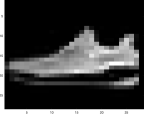
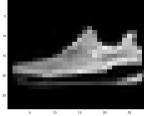
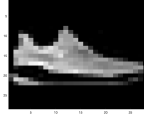
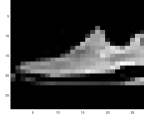
There are far too many combinations of layers, neurons, transfer functions, and training parameters to search for an optimal configuration methodically. Because of the time limitation, we decided to try a few reasonable combinations to find the one that works the best. We constrained the architectures in a few general ways. First, we limited the depth of our networks to 3-4 layers. This is slightly deeper than our best MNIST architecture, and a harder task like Fashion-MNIST would warrant more parameters in total. Second, each layer in the network has fewer neurons than the previous layer. The motivation for this funnel-like architecture is to help the layers learn incrementally more general features as they decrease in size. Most of our experimentation was with the transfer functions and the hyperparameters of learning rate, momentum, and weight decay. Cross-entropy loss was used when the output layer had softmax while MSE loss was used in all other cases. *Appendix 7.7 Loss Functions and their Derivatives* contains aforementioned functions.

## 2.3 Data Augmentation

The increase in performance of our model became very minimal over time. Rather than exploring further with different parameters with the same training data set, we decided to perform augmentation on the data set. Our model could benefit from having more data because it prevents over-fitting and would cover more examples from the entire distribution.

Our training data was augmented in 4 different ways as shown in Table 2. Each sample was rotated clock-wise and counter clock-wise by 2 degrees, flipped horizontally, and distorted with a set of x and y scale. Applying these augmentation on 50,000 samples, we gained 200,000 more training data.

Table 2. Example of Data Augmentations

Original	2° Clockwise	2° Counterclock.	Flipped	Distorted
				

An original image from Fashion-MNIST (left) compared with our augmentations of it. The differences between the augmentations and the original are barely perceivable by the human eye. However, they did result in a notable increase in accuracy after being added to the training set.

As the size of the training data quadrupled, the time spent in training a model significantly increased, taking as long as 8 hours. However, the increase in validation accuracy due to augmentation was undeniable. We were convinced that the augmentation of data will contribute a lot in improving the performance of our models. Hence, we started training all of our models with augmented data, sometimes even adding one more augmentation. Table 1 shows our experiment with data augmentation and its results in terms of accuracy. *Appendix 7.9 Augmented Datasets* demonstrates how the data was augmented using Python. *Appendix 7.13 Main Function - Augmented Data* demonstrates the driver code that builds and trains the model using the augmented data.

### 3 EXPERIMENTS - RESULTS

According to these experiments, using sigmoid for all hidden layers yields the best performance. Using Tanh worked extremely bad even though it looks a lot like sigmoid, and cursory reading states that these two functions should perform similarly. We do not know the reason for this poor performance. Using Relu worked surprisingly well as an output function, only slightly surpassed by softmax. We stuck to using momentum of 1 after attempting several different values of range  $[0, 1]$ . However, we acknowledge that momentum of 1 is equivalent to momentum of 0 with the weight update being offset by one batch. A lower learning rate seems to work best for 4-layer models as shown by the last line of the "50,000 train" section. Conversely, a higher learning rate seems to work better for 3-layer models as shown by the last layer of the "250,000 train" section. Small weight decay also improves performance slightly, as shown by the final two rows of the table. Data augmentation was effective at improving performance, increasing accuracy by 1.14% over models trained with original data. This was expected because more data gives the model more information with which to fit a function.

### 4 ENSEMBLE - METHODS

Our final attempt at improving performance was to aggregate our pile of trained models into a single, more intelligent model. The idea is to generate a prediction from many models, and weigh each model's prediction by how reliable it is at predicting a certain class. The reasoning is very intuitive. If model-1 predicts that the input is a shoe, but its classification performance on shoes is poor, its prediction is not to be trusted. Conversely, if model-2 predicts that the input is a dress, and its dress performance is very good, its prediction is deemed trustworthy. In practice, we used a class-wise performance vector over the validation set. This vector is the same shape as the prediction vectors, and each index holds the validation performance of its respective class. Mathematically, the ensemble is the weighted sum of each model's prediction multiplied by its respective performance vector. The ideal case is where each model has specialized in a unique subset of the classes so that they cover each other's shortcomings in the weighted sum.

In an ensemble  $E$  containing  $n$  models, the feedforward operation given an input vector  $x$  is defined as

$$E(x) = \text{hardmax} \left( \sum_{i=1}^n p_i \odot m_i(x) \right)$$

where  $m_i$  is the  $i$ -th model in the ensemble,  $p_i$  is the  $i$ -th model's class-wise performance vector over the validation set, and  $\odot$  is element-wise multiplication. *Appendix 7.3 Ensemble Class* demonstrates the ensemble model class.

#### 4.1 Choosing a Performance Metric

Table 3. Ensemble Accuracies Using Different Prediction Weights

Ensemble Config	Validation Accuracy %		
	Using Precision	Using Recall	Using (Precision $\odot$ Recall)
1	89.92	89.78	89.79
2	89.80	89.57	89.51
3	90.00	89.83	89.79
4	90.32	90.32	90.24
5	89.90	89.71	89.64
6	89.72	89.49	89.49

Validation accuracies for six random ensemble configurations, each tested with three different performance metrics. For every ensemble configuration, using precision results in the same or better accuracy than the other metrics.

There are several performance metrics to choose from, and it is not clear in theory which one would work the best. We can at least eliminate accuracy, because the vast majority of predictions and labels over any one class will be 0, i.e. negative. Since the class-wise prediction and label vectors are so sparse, the model will guess 0 most of the time and be correct, causing class-wise accuracy to be severely inflated. This leaves precision and recall, which solve this sparsity problem by only caring about positive predictions and labels.

We tested six different ensemble models, each containing a random and unique assortment of models. For each ensemble, we tested three different performance metrics: precision, recall, and the element-wise product of them. As shown by Table 3, ensembles using precision consistently outperform the same ensembles using the other two metrics. As a result, we exclusively used precision as the performance metric in every subsequent ensemble.

#### 4.2 Selecting Which Models to Ensemble

To find the best subset of models for the ensemble, we looked at each model's precision vector and chose the models with the highest precision in at least one class. Table 4 shows which models were chosen and the classes for which they had the highest precision.

*Appendix 7.4 Selecting Models for Ensemble* demonstrates the code used to analyse the precisions of each model and select an optimal set of models to ensemble. *Appendix 7.14 Main Function - Ensemble* shows the driver code used to assemble an ensemble of models and use it to generate a submission.

### 5 ENSEMBLE - RESULTS

We did see some gains in accuracy from the ensemble but they were minimal. The largest increase we saw was +0.2%, both in the validation accuracy as shown in Table 5 and the public Kaggle leaderboard. We came to a conclusion that the ensembling caused such small improvements because

Table 4. Classification Precision per Clothing Class

Model Name	Class Precisions									
	0	1	2	3	4	5	6	7	8	9
<b>HansModal</b>	<b>0.853</b>	0.986	0.827	<b>0.905</b>	<b>0.831</b>	0.965	<b>0.783</b>	0.956	0.964	0.962
<b>AdalynModelina</b>	0.84	0.986	<b>0.837</b>	0.891	0.825	0.969	0.77	0.951	0.97	0.966
<b>KateModelle</b>	0.84	<b>0.99</b>	0.789	0.886	0.79	0.961	0.783	<b>0.961</b>	0.966	0.951
aug-model-3-4-20-35	0.841	0.984	0.827	0.893	0.795	0.963	0.773	0.951	0.961	0.965
<b>aug-model-3-4-5-9</b>	0.824	0.983	0.793	0.893	0.81	<b>0.973</b>	0.776	0.959	<b>0.97</b>	0.95
model-3-3-5-14	0.845	0.978	0.813	0.878	0.808	0.97	0.743	0.954	0.955	0.964
<b>aug-model-3-7-3-11</b>	0.838	0.989	0.817	0.89	0.785	0.962	0.772	0.949	0.962	<b>0.967</b>
model-3-3-5-14	0.845	0.978	0.813	0.878	0.808	0.97	0.743	0.954	0.955	0.964

Classification precision for our eight best models per-class. The highest precision in each class is bolded, and the name of any model that achieved at least one highest precision is also bolded. Only these models that have at least one highest precision are used in the ensemble.

Table 5. Ensemble Accuracies After Selecting Models by Best Precision

Ensemble Config	Validation Accuracy %)
1	90.49
2	90.60

Validation accuracies for two ensemble configurations. For these ensembles, their model sets were chosen using the "best precision" strategy. Ensemble 1 was the first attempt using this strategy. Ensemble 2 was created using the same strategy after new models were introduced.

all of our models have an extremely similar distributions of competency. Table 4 shows that all the models have almost the same precisions to each other for every class, meaning they are smart and dumb in all the same areas. This is far from the ideal case for ensembling as described above, where each class would be covered by a model that excels at classifying it. There was not much to be gained from ensembling our particular set of models.

## 6 CONCLUSION

We were surprised how far the multilayer perceptron could be pushed on the Fashion-MNIST dataset. By far, the most improvement came from fine-tuning the MLP architecture and hyperparameters using the validation accuracy as a metric. We found that 3-4 layers worked well, along with a learning rate of 0.03-0.07, momentum of 1 (functionally equivalent to momentum of 0), and weight decay of -0.01. The best-performing function set proved to be sigmoid for all hidden layers, softmax for the output layer, and cross-entropy for the loss function, although relu on the output layer and MSE loss also worked extremely well together.

The next most significant technique was data augmentation, which resulted in a 1.14% accuracy increase from the best model trained on non-augmented data. This is an expected result of data augmentation because it helps diversify our data set while gaining more data.

The final and smallest increase in performance was achieved by carefully selecting a set of fully trained models and combining them into an ensemble. This technique caused a 0.2% increase in accuracy, both on the validation set and the Kaggle leaderboard when compared to the best individual model. Since all of our models had very similar distributions of class-wise performance, they could not complement each other very well in the ensemble's weighted sum. Nonetheless, their small differences were sufficient to achieve a notable increase in performance.

## 7 APPENDIX

### 7.1 Multilayer Perceptron Class

Listing 1. Multilayer Perceptron

```
1 classdef MultilayerPerceptron < handle
2
3     properties
4         layers
5         cost_func
6         d_cost_func
7     end
8
9     methods
10         function obj = MultilayerPerceptron(cost_func, d_cost_func)
11             obj.layers = [];
12             obj.cost_func = cost_func;
13             obj.d_cost_func = d_cost_func;
14         end
15
16         function add_layer(obj, layer)
17             obj.layers = [obj.layers layer];
18         end
19
20         function [losses, best_metric, all_metrics] = fit(obj, examples, labels,
21             epochs, batch_size, stop_buff, stop_thresh, test_data, test_labels,
22             metric_func, checkpoint_name)
23             for i = 1:size(obj.layers, 2)
24                 disp(size(obj.layers(i).W));
25             end
26
27             avg_losses = [];
28
29             batch_losses = zeros(batch_size);
30             epoch_losses = [];
31             temp_avg_losses = [];
32             metric = [];
33             best_metric = 0
34             disp([0.0000 metric_func(hardmax(obj.frozen_forward(test_data)),
35                 test_labels)]);
36
37             for e = 1:epochs
38                 perm = randperm(size(examples, 2));
39                 examples = examples(:, perm);
40                 labels = labels(:, perm);
41
42                 for b = 1:int32(ceil(size(examples, 2) / batch_size) + 1)
43                     % fprintf('batch %d / %d\n', b, int32(ceil(size(examples, 2) /
44                         batch_size) + 1));
45
46                     start_i = (b - 1) * (batch_size) + 1;
47                     end_i = min(b * batch_size, size(examples, 2));
48                     if start_i > end_i
49                         break
50                     end
51                 end
52             end
53         end
54     end
55 end
```

```

47         for i = start_i:end_i
48             a = obj.forward(examples(1:size(examples, 1), i));
49             batch_losses(i - start_i + 1) = obj.backward(a, labels(1:size(
                labels, 1), i));
50         end
51         obj.update((e-1)/epochs); % pass current progress
52         avg_losses = [avg_losses mean(batch_losses(1:(end_i - start_i + 1)
                ))];
53     end
54
55     metric = [metric metric_func(hardmax(obj.frozen_forward(test_data)),
        test_labels)];
56     disp([e metric(size(metric, 2))]);
57
58     if(metric(length(metric)) > best_metric)
59         best_metric = metric(length(metric));
60         obj.try_save_checkpoint(checkpoint_name, best_metric);
61     end
62
63     temp_avg_losses = avg_losses(:,:);
64     epoch_losses = [epoch_losses mean(temp_avg_losses)];
65     if(doEarlyStop(e, metric, stop_buff, stop_thresh))
66         break
67     end
68 end
69
70 losses = epoch_losses;
71 all_metrics = metric;
72 end
73
74 function try_save_checkpoint(obj, filename, metric)
75
76     % save current model
77     metric_str = num2str(metric);
78     model_timestamp = filename + "_METRIC_" + metric_str(3:size(metric_str, 2)
        ) + '.mat';
79     mlp = obj;
80     save(model_timestamp, "mlp");
81
82     % find and delete the other model
83     files = ls('models/*.mat');
84
85     filename = extractAfter(filename, "models/") + "_METRIC_";
86
87     for n = 1:height(files)
88
89         if contains(files(n,:), filename) & (strtrim(files(n,:)) ~=
            extractAfter(model_timestamp, "models/")) % to avoid deleting the
                current model
90             delete("models/" + files(n,:));
91             break
92         end
93     end
94 end

```

```

95
96     function out = forward(obj, vec)
97         for layer = obj.layers
98             vec = layer.forward(vec);
99         end
100         out = vec;
101     end
102
103     function out = frozen_forward(obj, vec)
104         for layer = obj.layers
105             vec = layer.frozen_forward(vec);
106         end
107         out = vec;
108     end
109
110     function loss = backward(obj, last_a, target)
111         last_n = obj.layers(size(obj.layers, 2)).n;
112         last_s = obj.layers(size(obj.layers, 2)).d_trans_func(last_n) .* obj.
            d_cost_func(last_a, target);
113
114         obj.layers(size(obj.layers, 2)).add_to_s(last_s);
115
116         for i = (size(obj.layers, 2) - 1):-1:1
117             obj.layers(i).backward(obj.layers(i+1).W, obj.layers(i+1).s)
118         end
119         loss = mean(obj.cost_func(last_a, target));
120     end
121
122     function update(obj, epoch_progress)
123         for i = 1:size(obj.layers)
124             obj.layers(i).update(epoch_progress);
125         end
126     end
127 end
128 end

```

## 7.2 Perceptron Layer Class

Listing 2. Perceptron Layer

```

1  classdef PerceptronLayer < handle
2      properties
3          trans_func % layer transfer function
4          d_trans_func % derivative of layer transfer function
5          batch_count % current count of how many examples have been seen in the current
                        batch
6          lr_max
7          lr_min
8          W % layer weights
9          b % layer biases
10         n % most recent net input, needed for backprop (backward)
11         a % most recent activation, needed for last layer sensitivity
12         s % most recent sensitivity, needed for backprop

```



```

13     p % most recent input vector, needed to compute each p's sensitivity (add_to_s
14     )
15     avg_s % sum of sensitivities in the current batch, needed for gradient descent
16     (update)
17     avg_sp % sum of the sensitivity * p-input' matrices in the current batch.
18     Averaged at update
19     last_W_update
20     last_b_update
21     momentum
22 end
23
24 methods
25
26     function obj = PerceptronLayer(arg1, arg2, trans_func, d_trans_func,
27         learn_rate_max, learn_rate_min, momentum, std)
28
29         obj.trans_func = trans_func;
30         obj.d_trans_func = d_trans_func;
31         obj.lr_max = learn_rate_max;
32         obj.lr_min = learn_rate_min;
33         obj.momentum = momentum;
34
35         % If both arguments are scalars, create random weight matrix and
36         % bias vector. Otherwise, use the provided weights and biases
37         if size(arg1) == [1 1] & size(arg2) == [1 1]
38             obj.W = normrnd(0, std, arg1, arg2);
39             obj.b = normrnd(0, std, arg1, 1);
40         else
41             obj.W = arg1;
42             obj.b = arg2;
43         end
44
45         obj.reset_for_next_batch();
46         obj.last_W_update = zeros(size(obj.W));
47         obj.last_b_update = zeros(size(obj.W, 1), 1);
48     end
49
50     function output = forward(obj, p)
51         obj.batch_count = obj.batch_count + 1;
52         obj.p = p;
53
54         obj.n = obj.W * p + obj.b;
55         obj.a = obj.trans_func(obj.n);
56         output = obj.a;
57     end
58
59     function output = frozen_forward(obj, p)
60         output = obj.trans_func(obj.W * p + obj.b);
61     end
62
63     function backward(obj, next_W, next_s)
64         obj.add_to_s(obj.d_trans_func(obj.n) .* (next_W' * next_s));
65     end
66
67     function update(obj, epoch_progress)
68         obj.avg_s = obj.avg_s / obj.batch_count;

```

```

63         obj.avg_sp = obj.avg_sp / obj.batch_count;
64
65         if obj.lr_min == obj.lr_max
66             lr = obj.lr_min;
67         else
68             lr = obj.lr_max - (obj.lr_max - obj.lr_min) * epoch_progress;
69         end
70
71         obj.W = obj.W + ((lr * (1 - obj.momentum) * obj.avg_sp) + (obj.momentum *
72             obj.last_W_update));
73         obj.b = obj.b + ((lr * (1 - obj.momentum) * obj.avg_s) + (obj.momentum *
74             obj.last_b_update));
75
76         obj.reset_for_next_batch();
77     end
78
79     function add_to_s(obj, s_vec)
80         obj.avg_s = obj.avg_s + s_vec;
81         obj.avg_sp = obj.avg_sp + (s_vec * obj.p');
82         obj.s = s_vec;
83     end
84
85     function reset_for_next_batch(obj)
86         obj.last_W_update = obj.avg_sp;
87         obj.last_b_update = obj.avg_s;
88
89         obj.avg_sp = zeros(size(obj.W));
90         obj.avg_s = zeros(size(obj.W, 1), 1);
91
92         obj.batch_count = 0;
93     end
94
95     function print(obj)
96         disp('Weights');
97         disp(obj.W);
98         disp('');
99
100         disp('Biases');
101         disp(obj.b);
102         disp('');
103     end
end

```

### 7.3 Ensemble Class

Listing 3. Ensemble Class

```

1 classdef Ensemble < handle
2
3     properties
4         model_list
5     end
6

```

```

7  methods
8      function obj = Ensemble()
9          obj.model_list = [];
10     end
11
12     function add_model(obj, model)
13         % add model to list
14         obj.model_list = [obj.model_list model];
15     end
16
17     function [losses, best_metrics, all_metrics] = fit(obj, examples, labels,
18         epochs, batch_size, stop_buff, stop_thresh, test_data, test_labels,
19         metric_func)
20         % fit all models in model_list one after another. Also, record
21         % loss, best_metric, and all_metrics for each model in a
22         % matrix, where each row is the result from one model.
23
24         losses = [];
25         best_metrics = [];
26         all_metrics = [];
27
28         for i = 1:size(obj.model_list, 2)
29             [curr_losses, curr_best_metrics, curr_all_metrics] = obj.model_list(i)
30                 .fit( ...
31                     examples, ...
32                     labels, ...
33                     epochs, ...
34                     batch_size, ...
35                     stop_buff, ...
36                     stop_thresh, ...
37                     test_data, ...
38                     test_labels, ...
39                     metric_func ...
40                 );
41             losses = [losses; curr_losses];
42             best_metrics = [best_metrics; curr_best_metrics];
43             all_metrics = [all_metrics; curr_all_metrics];
44         end
45
46         disp('FINAL INDIVIDUAL METRICS');
47         disp(best_metrics);
48         disp('FINAL ENSEMBLE METRIC');
49         disp(metric_func(hardmax(obj.frozen_forward(test_data, test_data,
50             test_labels)), test_labels));
51     end
52
53     function weighted_votes = frozen_forward(obj, vec, test_data, test_labels)
54         % perform the voting algorithm. The final votes are the average
55         % of the weighted sum of each model's predictions, where the weights are
56         % the
57         % model's (precision * recall) score per-class.
58
59         weighted_votes = zeros(size(test_labels));

```

```

56     for i = 1:size(obj.model_list, 2)
57         num_layers = size(obj.model_list(i).layers, 2);
58
59         % temporarily swap out the last layer transfer function for softmax
60         % temp_handle = obj.model_list(i).layers(num_layers).trans_func;
61         % obj.model_list(i).layers(num_layers).trans_func = @softmax;
62
63         % get predictions
64         preds = obj.model_list(i).frozen_forward(vec);
65
66         % revert last layer transfer function to original function
67         % obj.model_list(i).layers(num_layers).trans_func = temp_handle;
68
69         % compute model weights
70         confidence_weights = repmat(compute_prec_rec_weight(hardmax(preds),
71             test_labels), 1, size(vec, 2));
72
73         % add to the weighted sum
74         weighted_votes = weighted_votes + (confidence_weights .* preds);
75     end
76
77     % average the weighted sum
78     weighted_votes = weighted_votes * (1 / size(vec, 2));
79 end
80 end

```

## 7.4 Selecting Models for Ensemble

Listing 4. Function that selects the models with the best performance in each class

```

1  function precision_testing()
2      % make all folders visible to matlab
3      addpath('cost_functions');
4      addpath('d_cost_functions');
5      addpath('transfer_functions');
6      addpath('d_transfer_functions');
7      addpath('utility');
8      addpath('data');
9      addpath('nn_components');
10     addpath('augmented');
11
12     all_submission_data = readmatrix('test.csv'); % read all 10,000 submission
13     datapoints into matrix
14     submission_data = all_submission_data(:, 2:785)' * (1/255); % get rid of the
15     useless "id" column in the submission file
16
17     % split training and validation data
18     all_examples = readmatrix('train.csv');
19     all_labels = to_one_hot(all_examples(:, 2), 0, 9);
20     all_examples = all_examples(:, 3:786)' * (1/255);
21
22     % Training datapoints out of 60,000. The rest are used for validation
23     TRAIN_SIZE = 50000;

```

```

22 valid_data = all_examples(:, (TRAIN_SIZE + 1):60000);
23 valid_labels = all_labels(:, (TRAIN_SIZE + 1):60000);
24
25 % LOAD MODELS
26 % load mlp models from .mat files
27 model_files = [
28     "models/HansModal_2021_3_12_22_17_METRIC_904.mat"
29     "models/AdalynModelina_2021_3_12_1_15_METRIC_9014.mat"
30     "models/KateModelle_aug_2021_3_11_14_36_METRIC_8926.mat"
31     "models/aug_model_2021_3_4_20_35.mat"
32     "models/aug_model_2021_3_4_5_9.mat"
33     "models/model_2021_3_3_5_14.mat"
34     "models/aug_model_2021_3_7_3_11.mat"
35     "models/model_2021_3_2_14_27.mat"
36     "models/model_2021_3_2_17_49.mat"
37     "models/model_2021_3_3_5_14.mat"];
38
39 precisions = [];
40 best_model_per_class = [];
41 unique_models = [""];
42
43 for i = 1:length(model_files)
44     load(model_files(i), "mlp");
45     preds = hardmax(mlp.frozen_forward(valid_data));
46     precisions = [precisions compute_prec_rec_weight(preds, valid_labels)];
47 end
48 disp(precisions);
49
50 for i = 1:size(precisions, 1)
51     class_precs = precisions(i, 1:size(precisions, 2));
52     [maxval, maxi] = max(class_precs);
53     class_best = model_files(maxi);
54
55     best_model_per_class = [best_model_per_class; class_best];
56     if ~ismember(class_best, unique_models)
57         unique_models = [unique_models; class_best];
58     end
59 end
60 disp(best_model_per_class);
61 disp(unique_models);
62 end

```

Listing 5. Function that computes precision and recall

```

1 function weights = compute_prec_rec_weight(predictions, labels)
2     true_pos = sum(predictions & labels, 2);
3     precisions = true_pos ./ sum(predictions, 2); % TP / (TP + FP)
4     recalls = true_pos ./ sum(labels, 2); % TP / (TP + FN)
5
6     weights = precisions;
7 end

```

## 7.5 Accuracy

Listing 6. Accuracy Function

```

1 function a = accuracy(predictions, labels)
2     count = 0;
3     label_size = size(labels);
4
5     for i = 1:label_size(2)
6         if predictions(1:label_size(1), i) == labels(1:label_size(1), i)
7             count = count + 1;
8         end
9     end
10    a = count / label_size(2);
11 end

```

## 7.6 Transfer Functions and their Derivatives

Listing 7. Linear

```

1 function l = linear(x)
2     l = x;
3 end

```

Listing 8. Linear Derivative

```

1 function l = d_linear(x)
2     l = ones(size(x));
3 end

```

Listing 9. Tanh

```

1 function t = my_tanh(x)
2     t = (exp(x) - exp(-x)) .* ((exp(x) + exp(-x)) .^ (-1));
3 end

```

Listing 10. Tanh Derivative

```

1 function t = d_my_tanh(x)
2     t = 1 - (my_tanh(x) .^ 2);
3 end

```

Listing 11. ReLU

```

1 function r = relu(x)
2     x(x < 0) = 0;
3     r = x;
4 end

```

Listing 12. ReLU Derivative

```

1 function r = d_relu(x)
2     x(x < 0) = 0;
3     x(x >= 0) = 1;
4     r = x;

```

```
5 end
```

Listing 13. Sigmoid

```
1 function s = sigmoid(x)
2     s = (1 + exp(-x)) .^ (-1);
3 end
```

Listing 14. Sigmoid Derivative

```
1 function s = d_sigmoid(x)
2     s = sigmoid(x) .* (1 - sigmoid(x));
3 end
```

Listing 15. Softmax

```
1 function s = softmax(x)
2     exp_of_x = exp(x);
3     s = exp_of_x ./ sum(exp_of_x, 1);
4 end
```

Listing 16. Softmax Derivative

```
1 function s = d_softmax(x)
2     % softmax(x) is always used as softmax(sigmoid(x)), so we apply the chain rule
   here:
3     % d_softmax(sigmoid(x)) * d_sigmoid(x)
4     s = softmax(x) .* (1 - softmax(x));
5 end
```

## 7.7 Loss Functions and their Derivatives

Listing 17. Mean Squared Error

```
1 function error = squared_error(a, t)
2     e = a - t;
3     error = e .* e;
4 end
```

Listing 18. Mean Squared Error Derivative

```
1 function error = d_squared_error(a, t)
2     error = t - a;
3 end
```

Listing 19. Cross-Entropy

```
1 function loss = cross_entropy(a, t)
2     loss = -sum(t .* log(a));
3 end
```

Listing 20. Cross-Entropy Derivative (softmax derivative included)

```

1 function error = d_cross_entropy(a, t)
2     % error = -(t ./ a) + ((1-t) ./ (1-a));
3     error = t - a;
4 end

```

## 7.8 Hardmax Function

Listing 21. Hardmax

```

1 function out = hardmax(x)
2     for i = 1:size(x, 2)
3         [max_num, max_index] = max(x(1:size(x, 1), i));
4         x(1:size(x, 1), i) = zeros(size(x, 1), 1);
5         x(max_index, i) = 1;
6     end
7     out = x;
8 end

```

## 7.9 Generating Augmented Datasets

Listing 22. Dataset Augmentation Script (Python)

```

1 """
2 Functions
3 - rotate, blur, add_salt_pepper_noise, distort, translate, scale, flip
4 - augment: performs some or all of those augmentation for each data sample in train
   set
5
6 Instructions
7 1) In 'uwb-css-485-winter-2021' directory, make duplicate of 'train.csv' and name it '
   train_for_aug.csv'
8 2) Make following changes to 'train_for_aug.csv'
9     - delete the first row ("id", "label", "pixel1", "pixel2", ...)
10    - delete the validation data (rows from 50,001 - 60,000)
11 3) In the same depth as the 'utility' folder, make a directory called 'augmented'
12 4) Currently this code performs 4 different augmentation. if you want, make changes to
   augment() function
13 5) Run this code and the output will be saved in 'augmented' directory
14 6) When you train your model on MATLAB, train with main_agumented() instead of main()
15
16 """
17 import pandas as pd
18 import numpy as np
19 from numpy import genfromtxt
20 from IPython.display import Image
21 from PIL import Image as PILImage
22 import scipy.ndimage
23 import cv2
24 import random
25 import csv
26 from csv import reader, writer
27 from datetime import datetime
28

```



```

29 import gc
30
31 def blur(sample, sigma=0.58):
32     return scipy.ndimage.filters.gaussian_filter(sample, sigma=sigma)
33
34 def add_salt_pepper_noise(sample, prob=0.05):
35     output = np.zeros(sample.shape,np.uint8)
36     thres = 1 - prob
37     for i in range(sample.shape[0]):
38         for j in range(sample.shape[1]):
39             rdn = random.random()
40             if rdn < prob:
41                 output[i][j] = 0
42             elif rdn > thres:
43                 output[i][j] = 255
44             else:
45                 output[i][j] = sample[i][j]
46     return output
47
48 def distort(img, orientation='horizontal', func=np.sin, x_scale=0.05, y_scale=5):
49     assert orientation[:3] in ['hor', 'ver'], "dist_orient should be 'horizontal'|"
50         vertical'"
51     assert func in [np.sin, np.cos], "supported functions are np.sin and np.cos"
52     assert 0.00 <= x_scale <= 0.1, "x_scale should be in [0.0, 0.1]"
53     assert 0 <= y_scale <= min(img.shape[0], img.shape[1]), "y_scale should be less
54         then image size"
55     img_dist = img.copy()
56
57     def shift(x):
58         return int(y_scale * func(np.pi * x * x_scale))
59
60     for _ in range(3):
61         for i in range(img.shape[orientation.startswith('ver')]):
62             if orientation.startswith('ver'):
63                 img_dist[:, i] = np.roll(img[:, i], shift(i))
64             else:
65                 img_dist[i, :] = np.roll(img[i, :], shift(i))
66
67     return img_dist
68
69 def translate(sample, shift=1, direction = 1):
70     # direction 1,2,3,4 correstponds to
71     # left, right, upward, downward
72
73     if direction < 3:
74         shifted_area = np.zeros((28, shift))
75         ax = 1
76     else:
77         shifted_area = np.zeros((shift, 28))
78         ax = 0
79
80     if direction == 1:
81         sample = sample[:, shift:]
82     elif direction == 2:

```

```

81     sample = sample[:, :-shift]
82 elif direction == 3:
83     sample = sample[shift:, :]
84 else:
85     sample = sample[:, :-shift, :]
86
87 if direction % 2 == 0: # downward and right
88     shifted = np.concatenate((shifted_area, sample), axis=ax)
89 else:
90     shifted = np.concatenate((sample, shifted_area), axis=ax)
91
92 return shifted
93
94 def scale(sample, dsize=(28,28), resultsize=(28,28), interpolation=cv2.INTER_CUBIC):
95     res = cv2.resize(sample, dsize=dsize, interpolation=cv2.INTER_CUBIC)
96
97     if dsize[0] > resultsize[0]: # cut from each side almost equally
98         diff = dsize[0] - resultsize[0]
99         res = res[:, diff//2 : -(diff)//2]
100
101     if dsize[1] > resultsize[1]:
102         diff = dsize[1] - resultsize[1]
103         res = res[diff//2 : -(diff)//2, :]
104
105     return res
106
107 def rotate(sample, angle=2):
108     return scipy.ndimage.rotate(sample, angle, reshape=False)
109
110 def flip(sample):
111     return np.fliplr(sample)
112
113 def augment(filename):
114     t = datetime.now().strftime("%Y_%m_%d_%H_%M_%S")
115     count = 0
116
117     with open(f'../augmented/augmented_train_{t}.csv', 'a', newline='') as f1:
118         csvwriter1 = writer(f1)
119         with open(f'../augmented/augmented_label_{t}.csv', 'a', newline='') as f2:
120             csvwriter2 = writer(f2)
121
122         with open(filename, 'r') as read_obj:
123             csv_reader = reader(read_obj)
124
125             for row in csv_reader:
126                 count += 1
127                 if count % 1000 == 0: print(f'Progress: {count}/50000')
128                 if count % 500 == 0: gc.collect()
129
130                 arr = np.asarray(row).astype(float)
131                 label = arr[1]
132                 arr = arr[2:]
133
134                 sample = arr.reshape(28,28)

```

```

135
136     """
137     !!! Make changes here if you want more or less augmentation !!!
138     Don't forget to .ravel()
139     Don't forget to change the range(n) in the for loop for
        writing labels
140     """
141     # writing origianl and augmented images
142     csvwriter1.writerow([sample.ravel(),\
143                         rotate(sample,-2).ravel(),\
144                         rotate(sample, 2).ravel(),\
145                         flip(sample).ravel(),\
146                         distort(sample, x_scale=0.03, y_scale=2).ravel
                            ( )])
147
148     # writing labels
149     for _ in range(5):
150         csvwriter2.writerow([label])
151
152     return t
153
154 def main():
155
156     t = augment("../uwb-css-485-winter-2021/train_for_aug.csv") # this file should not
        contain the first row and the rows from 50,001 – 60,000
157
158     print("\nFiles generated: ")
159     print(f'augmented/augmented_train_{t}.csv')
160     print(f'augmented/augmented_label_{t}.csv')
161
162 if __name__ == "__main__":
163     main()

```

## 7.10 Early Stopping

Listing 23. Function that decides whether to stop early

```

1 function do_stop = doEarlyStop(ep, values, early_stop_buff_size, early_stop_threshold)
2     % return true to do an early stop IF:
3     % the current error is 0
4     % OR BOTH the past <buff_size> errors have a standard deviation below
5     %         a threshold AND there have been at least <buff_size> epochs
6     if(ep >= early_stop_buff_size)
7         disp(mean_trend(values( (size(values, 2) – early_stop_buff_size + 1):size(
            values, 2) )))
8     end
9     do_stop = (ep >= early_stop_buff_size ...
10         && mean_trend(values( (size(values, 2) – early_stop_buff_size + 1):size(
            values, 2) )) <= early_stop_threshold);
11 end

```

Listing 24. Function that returns the average change in validation metric over N epochs

```

1 function t = mean_trend(vec)
2     if(size(vec, 2) == 1)

```

```

3     t = 0;
4     else
5         t = (vec(size(vec, 2)) - vec(1)) / ((size(vec, 2) - 1));
6     end
7 end

```

## 7.11 One-Hot Encoding and Decoding

Listing 25. One-Hot Encoding - Converts a list of ints to a matrix of one-hot column vectors

```

1 function one_hot = to_one_hot(colwise_nums, min, max)
2     one_hot = (colwise_nums == min:max)';
3 end

```

Listing 26. One-Hot Decoding - Converts a matrix of one-hot column vectors to a list of ints

```

1 function nums = one_hot_to_int(colwise_onehot, min, max)
2     nums = colwise_onehot * (min:max)';
3 end

```

## 7.12 Main Function - Standard

Listing 27. Driver code that initializes data builds and trains model and generates a submission file

```

1 function main()
2     % make all folders fisible to matlab
3     addpath('cost_functions');
4     addpath('d_cost_functions');
5     addpath('transfer_functions');
6     addpath('d_transfer_functions');
7     addpath('utility');
8     addpath('data');
9     addpath('nn_components');
10
11     % Training datapoints out of 60,000. The rest are used for validation
12     TRAIN_SIZE = 50000;
13
14     % READ ALL DATA
15
16     all_data = readmatrix('train.csv'); % read all 60,000 labeled datapoints and
17         labels into matrix
18     all_submission_data = readmatrix('test.csv'); % read all 10,000 submission
19         datapoints into matrix
20     submission_data = all_submission_data(:, 2:785)' * (1/255); % normalize, and get
21         rid of the useles "id" column in the submission file
22
23
24     all_examples = all_data(:, 3:786)' * (1/255); % normalize datapoints
25     disp(all_examples(1:20, 1:10));
26     all_labels = to_one_hot(all_data(:, 2), 0, 9); % convert labels (0-9) to one-hot
27         vectors
28
29     % split training and validation data
30     train_data = all_examples(:, 1:TRAIN_SIZE);
31     train_labels = all_labels(:, 1:TRAIN_SIZE);
32     test_data = all_examples(:, (TRAIN_SIZE + 1):60000);

```

```

28     test_labels = all_labels(:, (TRAIN_SIZE + 1):60000);
29
30     % hyperparameters
31     epochs = 1000;
32     stop_buff = 1;
33     stop_thresh = -1;
34     std = 0.4;
35     batch_size = 32;
36     momentum = 0;
37     lr = 0.05;
38
39     % build model
40     mlp = MultilayerPerceptron(@cross_entropy, @d_cross_entropy);
41
42     mlp.add_layer(PerceptronLayer(256, 784, @my_tanh, @d_my_tanh, lr, momentum, std));
43     mlp.add_layer(PerceptronLayer(10, 256, @softmax, @d_softmax, lr, momentum, std));
44
45     % train the model
46     [losses, acc, acc_list] = mlp.fit( ...
47         train_data, ...
48         train_labels, ...
49         epochs, ...
50         batch_size, ...
51         stop_buff, ...
52         stop_thresh, ...
53         test_data, ...
54         test_labels, ...
55         @accuracy ...
56     );
57     disp(losses);
58     disp(acc_list);
59
60     % predict labels for submission data
61     final_preds = one_hot_to_int(hardmax(mlp.frozen_forward(submission_data)), 0, 9);
62
63     % format matrix for submission
64     submission_matrix = [(60001:70000)' final_preds];
65
66     % write submission matrix to data/SUBMISSION.csv
67     writematrix(submission_matrix, 'data/SUBMISSION.csv');
68 end

```

### 7.13 Main Function - Augmented Data

Listing 28. Driver code that initializes the augmented data builds and trains model and generates a submission

```

1 % !!! IMPORTNAT !!!
2
3 % this file saves the model it trained to a directory called 'models'
4 % please make sure you have a directory 'models' before running this code
5
6 function main_augmented()
7     % make all folders fisible to matlab
8     addpath('cost_functions');

```

```

9     addpath('d_cost_functions');
10    addpath('transfer_functions');
11    addpath('d_transfer_functions');
12    addpath('utility');
13    addpath('data');
14    addpath('nn_components');
15    addpath('augmented');
16
17    % read data — CHANGE THE FILE TO YOUR AUGMENTED DATASET
18
19    aug_data = readmatrix('augmented/small_augmented_train.csv');
20    aug_label = readmatrix('augmented/small_agumented_label.csv');
21
22    % aug_data = readmatrix('augmented/augmented_train_2021_03_12_15_15_00.csv');
23    % aug_label = readmatrix('augmented/augmented_label_2021_03_12_15_15_00.csv');
24    disp(size(aug_label));
25
26    all_submission_data = readmatrix('test.csv'); % read all 10,000 submission
           datapoints into matrix
27    submission_data = all_submission_data(:, 2:785)' * (1/255); % get rid of the
           useless "id" column in the submission file
28
29    train_data = aug_data' * (1/255); % normalize datapoints
30    train_labels = to_one_hot(aug_label, 0, 9); % convert labels (0–9) to one-hot
           vectors
31
32    % split training and validation data
33    all_examples = readmatrix('train.csv');
34    all_labels = to_one_hot(all_examples(:, 2), 0, 9);
35    all_examples = all_examples(:, 3:786)' * (1/255);
36
37
38    % Training datapoints out of 60,000. The rest are used for validation
39    TRAIN_SIZE = 50000;
40    valid_data = all_examples(:, (TRAIN_SIZE + 1):60000);
41    valid_labels = all_labels(:, (TRAIN_SIZE + 1):60000);
42
43    % hyperparameters
44    epochs = 20;
45    stop_buff = 1;
46    stop_thresh = -1;
47    std = 0.4;
48    batch_size = 24; %500;
49    momentum = 1.1;
50    lr_max = 1;
51    lr_min = 0.02;
52
53    % build model
54    mlp = MultilayerPerceptron(@cross_entropy, @d_cross_entropy);
55
56    mlp.add_layer(PerceptronLayer(350, 784, @sigmoid, @d_sigmoid, lr_max, lr_min,
           momentum, std));
57    mlp.add_layer(PerceptronLayer(170, 350, @sigmoid, @d_sigmoid, lr_max, lr_min,
           momentum, std));

```

```

58     mlp.add_layer(PerceptronLayer(10, 170, @relu, @d_relu, lr_max, lr_min, momentum,
59         std));
60
61     % change each run to make identifying models easier
62     model_name = "HansModal";
63
64     t = datetime('now');
65     model_timestamp = "models/" + model_name + "_" + year(t) + '_' + month(t) + '_' +
        day(t) + '_' + hour(t) + '_' + minute(t);
66
67     % train the model
68     [losses, acc, acc_list] = mlp.fit( ...
69         train_data, ...
70         train_labels, ...
71         epochs, ...
72         batch_size, ...
73         stop_buff, ...
74         stop_thresh, ...
75         valid_data, ...
76         valid_labels, ...
77         @accuracy, ...
78         model_timestamp ...
79     );
80     disp(" ===== losses ===== " );
81     disp(losses);
82     disp(" ===== acc_list ===== " );
83     disp(acc_list);
84
85     load("models/HansModal_2021_3_12_22_17_METRIC_904.mat", "mlp");
86     % predict labels for submission data
87     final_preds = one_hot_to_int(hardmax(mlp.frozen_forward(submission_data)), 0, 9);
88
89     % format matrix for submission
90     submission_matrix = [(60001:70000)' final_preds];
91
92     % write submission matrix to data/SUBMISSION.csv
93     submission_timestamp = "data/" + model_name + "_submission_" + year(t) + '_' +
        month(t) + '_' + day(t) + '_' + hour(t) + '_' + minute(t) + '.csv';
94     writematrix(submission_matrix, submission_timestamp);
95
96     % save the model
97     % model_timestamp = "models/models/" + model_name + "_" + year(t) + '_' + month(t)
        + '_' + day(t) + '_' + hour(t) + '_' + minute(t) + '.mat';
98     % save(model_timestamp, "mlp");
99 end

```

## 7.14 Main Function - Ensemble

Listing 29. Driver code that initializes data builds an Ensemble and uses it to generate a submission

```

1 function ensemble_main()
2     % make all folders fisible to matlab
3     addpath('cost_functions');

```

```

4  addpath('d_cost_functions');
5  addpath('transfer_functions');
6  addpath('d_transfer_functions');
7  addpath('utility');
8  addpath('data');
9  addpath('nn_components');
10 addpath('augmented');
11
12 all_submission_data = readmatrix('test.csv'); % read all 10,000 submission
      datapoints into matrix
13 submission_data = all_submission_data(:, 2:785)' * (1/255); % get rid of the
      useles "id" column in the submission file
14
15 % split training and validation data
16 all_examples = readmatrix('train.csv');
17 all_labels = to_one_hot(all_examples(:, 2), 0, 9);
18 all_examples = all_examples(:, 3:786)' * (1/255);
19
20 % Training datapoints out of 60,000. The rest are used for validation
21 TRAIN_SIZE = 50000;
22 valid_data = all_examples(:, (TRAIN_SIZE + 1):60000);
23 valid_labels = all_labels(:, (TRAIN_SIZE + 1):60000);
24
25
26 ensemble = Ensemble();
27
28 % LOAD MODELS
29 % load mlp models from .mat files
30 load("models/aug_model_2021_3_7_3_11.mat", "mlp");
31 ensemble.add_model(mlp);
32
33 load("models/KateModelle_aug_2021_3_11_14_36_METRIC_8926.mat", "mlp");
34 ensemble.add_model(mlp);
35
36 load("models/HansModal_2021_3_12_22_17_METRIC_904.mat", "mlp");
37 ensemble.add_model(mlp);
38
39 %   load("models/model_2021_3_3_5_14.mat", "mlp");
40 %   ensemble.add_model(mlp);
41
42 load("models/aug_model_2021_3_4_5_9.mat", "mlp");
43 ensemble.add_model(mlp);
44
45 load("models/AdalynModelina_2021_3_12_1_15_METRIC_9014.mat", "mlp");
46 ensemble.add_model(mlp);
47
48 disp(accuracy(hardmax(ensemble.frozen_forward(valid_data, valid_data, valid_labels
      )), valid_labels));
49
50 % predict labels for submission data
51 final_preds = one_hot_to_int(hardmax(ensemble.frozen_forward(submission_data,
      valid_data, valid_labels)), 0, 9);
52
53 % format matrix for submission

```



```
54 submission_matrix = [(60001:70000)' final_preds];
55
56 % write submission matrix
57 t = datetime('now');
58 submission_timestamp = "data/ensemble_submission_" + year(t) + '_' + month(t) + '_'
    ' + day(t) + '_' + hour(t) + '_' + minute(t) + '.csv';
59 writematrix(submission_matrix, submission_timestamp);
60
61 end
```